NLP Practice Notes — Week 1

Tokenization converts raw text into tokens. Sentence-level tokenization respects punctuation such as periods. Word-level tokenization further splits sentences.

Subword schemes such as BPE and WordPiece keep the vocabulary bounded while still representing rare words.

Week 2 - Metrics

BLEU scores precision-oriented n-gram overlap for machine translation, while ROUGE uses recall for summaries. Perplexity is the exponentiated cross-entropy per token.

Lower perplexity means the language model is less surprised by test data.